

Ryan.Kelly@harrison.ai

@rfkelly

<https://rfk.id.au/>



harrison.ai




harrison.ai

Indexing Petabytes of Data with Rust and AWS

10th October 2022




Amazon S3



A few PB of de-identified
medical image data in
millions tar archives





Amazon S3



A few PB of de-identified
medical image data in
millions tar archives




Amazon S3



A few PB of de-identified
medical image data in
millions tar archives



Amazon S3




Amazon Athena





A few PB of de-identified
medical image data in
millions tar archives



Amazon S3



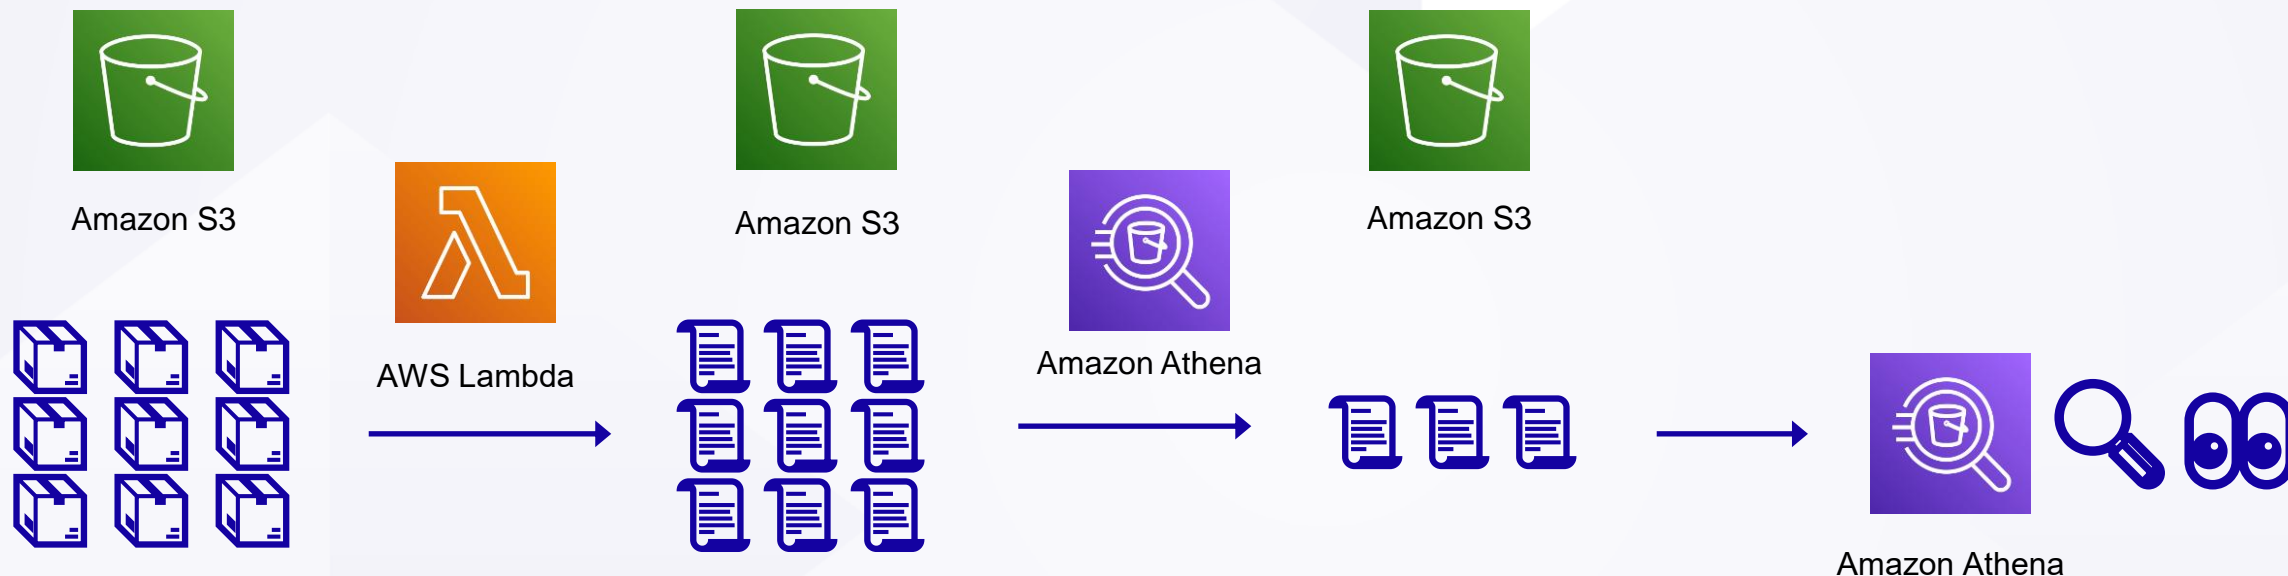
Amazon S3



Amazon Athena

A few PB of de-identified
medical image data in
millions tar archives


.parquet files
summarizing
available data




A few PB of de-identified
medical image data in
millions tar archives

.jsonl files listing
contents of each
archive

.parquet files
summarizing
available data




Amazon S3



A few PB of de-identified
medical image data in
millions tar archives




AWS Lambda




Amazon S3



.jsonl files listing
contents of each
archive



Amazon S3



A few PB of de-identified
medical image data in
millions tar archives



AWS Lambda



?



Amazon S3



.jsonl files listing
contents of each
archive

A Python Baseline

```
import json
import tarfile

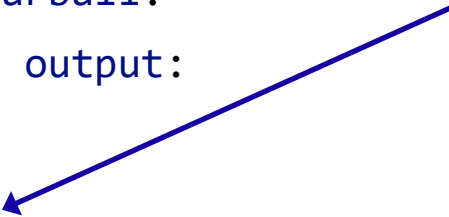
def index_tarball(input_path, output_path):
    with tarfile.open(input_path) as tarball:
        with open(output_path, "w") as output:
            for member in tarball:
                row = json.dumps({
                    "archive": input_path,
                    "filename": member.name,
                    "size": member.size
                })
                output.write(row)
                output.write("\n")
```

A Python Baseline

```
import json
import tarfile

def index_tarball(input_path, output_path):
    with tarfile.open(input_path) as tarball:
        with open(output_path, "w") as output:
            for member in tarball:
                row = json.dumps({
                    "archive": input_path,
                    "filename": member.name,
                    "size": member.size
                })
                output.write(row)
                output.write("\n")
```

IRL, this would be rich metadata about the file, e.g. opaque patient and study identifiers



A Basic Rust Equivalent

```
[dependencies]
anyhow = "1"
tar = "0.4"
serde = { version = "1", features=["derive"] }
serde_json = "1"
```

A Basic Rust Equivalent

```
use anyhow::{Context, Result};
use serde::Serialize;
use std::fs::File;
use std::io::prelude::*;

#[derive(Serialize)]
struct IndexEntry {
    archive: String,
    filename: String,
    size: u64,
}
```

A Basic Rust Equivalent

```
fn index_tarball(input_path: &str, output_path: &str) -> Result<()> {  
    let mut tarball = tar::Archive::new(File::open(input_path)?);  
    let mut output = File::create(output_path)?;  
    for entry in tarball.entries()? {  
        let entry = entry?;  
        let row = serde_json::to_string(&IndexEntry {  
            archive: input_path.into(),  
            filename: entry.path()?.to_str().context("non-utf8 path")?.into(),  
            size: entry.size(),  
        })?;  
        writeln!(output, "{}", row)?;  
    }  
    Ok(())  
}
```

A Basic Rust Equivalent

```
fn index_tarball(input_path: &str, output_path: &str) -> Result<()> {  
    let mut tarball = tar::Archive::new(File::open(input_path)?);  
    let mut output = File::create(output_path)?;  
    for entry in tarball.entries()? {  
        let entry = entry?;  
        let row = serde_json::to_string(&IndexEntry {  
            archive: input_path.into(),  
            filename: entry.path()?.to_str().context("non-utf8 path")?.into(),  
            size: entry.size(),  
        })?;  
        writeln!(output, "{}", row)?;  
    }  
    Ok(())  
}
```

Rust won't let a potential error pass
unhandled



Gut-check performance comparison

```
$ hyperfine "python ./src/0_naive_python/index-tarballs.py ./input"
```

```
Benchmark 1: python ./src/0_naive_python/index-tarballs.py ./input
```

```
Time (mean  $\pm$   $\sigma$ ):      345.7 ms  $\pm$  10.3 ms    [User: 328.8 ms, System: 16.9 ms]
```

```
Range (min ... max):    334.4 ms ... 360.3 ms    10 runs
```

Gut-check performance comparison

```
$ hyperfine "python ./src/0_naive_python/index-tarballs.py ./input"
```

```
Benchmark 1: python ./src/0_naive_python/index-tarballs.py ./input
```

```
Time (mean ±  $\sigma$ ):      345.7 ms ± 10.3 ms    [User: 328.8 ms, System: 16.9 ms]
```

```
Range (min ... max):    334.4 ms ... 360.3 ms    10 runs
```

```
$ hyperfine "cargo run --bin index-tarballs-1 ./input"
```

```
Benchmark 1: cargo run --bin index-tarballs-1
```

```
Time (mean ±  $\sigma$ ):      658.9 ms ± 11.6 ms    [User: 387.9 ms, System: 231.4 ms]
```

```
Range (min ... max):    644.4 ms ... 676.8 ms    10 runs
```

Gut-check performance comparison

```
$ cargo build --release
```

Gut-check performance comparison

```
$ cargo build --release
```

```
$ hyperfine "python ./src/0_naive_python/index-tarballs.py ./input"
```

```
Benchmark 1: python ./src/0_naive_python/index-tarballs.py ./input
```

```
Time (mean ± σ):      330.1 ms ±   8.4 ms    [User: 314.8 ms, System: 15.1 ms]
```

```
Range (min ... max):  321.1 ms ... 342.6 ms    10 runs
```

Gut-check performance comparison

```
$ cargo build --release
```

```
$ hyperfine "python ./src/0_naive_python/index-tarballs.py ./input"
```

```
Benchmark 1: python ./src/0_naive_python/index-tarballs.py ./input
```

```
Time (mean  $\pm$   $\sigma$ ):      330.1 ms  $\pm$  8.4 ms    [User: 314.8 ms, System: 15.1 ms]
```

```
Range (min ... max):    321.1 ms ... 342.6 ms    10 runs
```

```
$ hyperfine "./target/release/index-tarballs-1 ./input"
```

```
Benchmark 1: ./target/release/index-tarballs-1 ./input
```

```
Time (mean  $\pm$   $\sigma$ ):      215.7 ms  $\pm$  6.4 ms    [User: 29.6 ms, System: 186.1 ms]
```

```
Range (min ... max):    206.7 ms ... 225.4 ms    14 runs
```

Some Light Optimisation

```
fn index_tarball(input_path: &str, output_path: &str) -> Result<()> {  
    let mut tarball = tar::Archive::new(File::open(input_path)?);  
    let mut output = File::create(output_path)?;  
    for entry in tarball.entries()? {  
        let entry = entry?;  
        let row = serde_json::to_string(&IndexEntry {  
            archive: input_path.into(),  
            filename: entry.path()?.to_str().context("non-utf8 path")?.into(),  
            size: entry.size(),  
        })?;  
        writeln!(output, "{}", row)?;  
    }  
    Ok(())  
}
```

Some Light Optimisation

```
fn index_tarball(input_path: &str, output_path: &str) -> Result<()> {  
    let mut tarball = tar::Archive::new(File::open(input_path)?);  
    let mut output = File::create(output_path)?;  
    for entry in tarball.entries()? {  
        let entry = entry?;  
        let row = serde_json::to_string(&IndexEntry {  
            archive: input_path.into(),  
            filename: entry.path()?.to_str().context("non-utf8 path")?.into(),  
            size: entry.size(),  
        })?;  
        writeln!(output, "{}", row)?;  
    }  
    Ok(())  
}
```

Unnecessary string copies



Some Light Optimisation

```
fn index_tarball(input_path: &str, output_path: &str) -> Result<()> {  
    let mut tarball = tar::Archive::new(File::open(input_path)?);  
    let mut output = File::create(output_path)?;  
    for entry in tarball.entries()? {  
        let entry = entry?;  
        let row = serde_json::to_string(&IndexEntry {  
            archive: input_path,  
            filename: entry.path()?.to_str().context("non-utf8 path"?),  
            size: entry.size(),  
        })?;  
        writeln!(output, "{}", row)?;  
    }  
    Ok(())  
}
```


Some Light Optimisation

```
fn index_tarball(input_path: &str, output_path: &str) -> Result<()> {  
    let mut tarball = tar::Archive::new(File::open(input_path)?);  
    let mut output = File::create(output_path)?;  
    for entry in tarball.entries()? {  
        let entry = entry?;  
        let row = serde_json::to_string(&IndexEntry {  
            archive: input_path,  
            filename: entry.path()?.to_str().context("non-utf8 path"?),  
            size: entry.size(),  
        })?;  
        writeln!(output, "{}", row)?;  
    }  
    Ok(())  
}
```

← Unnecessary intermediate string

Some Light Optimisation

```
fn index_tarball(input_path: &str, output_path: &str) -> Result<()> {  
    let mut tarball = tar::Archive::new(File::open(input_path)?);  
    let mut output = File::create(output_path)?;  
    for entry in tarball.entries()? {  
        let entry = entry?;  
        serde_json::to_writer(&mut output, &IndexEntry {  
            archive: input_path,  
            filename: entry.path()?.to_str().context("non-utf8 path"?),  
            size: entry.size(),  
        })?;  
        writeln!(output)?;  
    }  
    Ok(())  
}
```

Some Light Optimisation

```
fn index_tarball(input_path: &str, output_path: &str) -> Result<()> {  
    let mut tarball = tar::Archive::new(File::open(input_path)?);  
    let mut output = File::create(output_path)?;  
    for entry in tarball.entries()? {  
        let entry = entry?;  
        serde_json::to_writer(&mut output, &IndexEntry {  
            archive: input_path,  
            filename: entry.path()?.to_str().context("non-utf8 path"?),  
            size: entry.size(),  
        })?;  
        writeln!(output)?;  
    }  
    Ok(())  
}
```

←← Lots of small reads/writes

Some Light Optimisation

```
fn index_tarball(input_path: &str, output_path: &str) -> Result<()> {  
    let mut tarball = tar::Archive::new(BufReader::new(File::open(input_path)?));  
    let mut output = BufWriter::new(File::create(output_path)?);  
    for entry in tarball.entries()? {  
        let entry = entry?;  
        serde_json::to_writer(&mut output, &IndexEntry {  
            archive: input_path,  
            filename: entry.path()?.to_str().context("non-utf8 path")?,  
            size: entry.size(),  
        })?;  
        writeln!(output)?;  
    }  
    Ok(())  
}
```

Gut-check performance comparison

```
$ hyperfine "python ./src/0_naive_python/index-tarballs.py ./input"
```

```
Benchmark 1: python ./src/0_naive_python/index-tarballs.py ./input
```

```
Time (mean ± σ):      330.1 ms ±   8.4 ms    [User: 314.8 ms, System: 15.1 ms]
```

```
Range (min ... max):  321.1 ms ... 342.6 ms    10 runs
```

```
$ hyperfine "./target/release/index-tarballs-2 ./input"
```

```
Benchmark 1: ./target/release/index-tarballs-2 ./input
```

```
Time (mean ± σ):      130.5 ms ±   5.8 ms    [User: 13.0 ms, System: 117.5 ms]
```

```
Range (min ... max):  123.6 ms ... 144.6 ms    20 runs
```

So anyway...AWS?

A Python Baseline

```
import boto3

def index_tarball(s3client, bucket, input_key, output_key):
    input = s3client.get_object(Bucket=bucket, Key=input_key)["Body"]

    output = BytesIO()
    with tarfile.open(fileobj=input, mode="r|") as tarball:
        for member in tarball:
            row = json.dumps(
                {"archive": input_key, "filename": member.name, "size": member.size}
            )
            output.write(row.encode("utf-8"))
            output.write(b"\n")

    output.seek(0)
    s3client.put_object(
        Bucket=bucket,
        Key=output_key,
        Body=output,
    )
```

A Rust Equivalent

```
[dependencies]
anyhow = "1"
async-tar = "0.4"
aws-config = "0.49"
aws-sdk-s3 = "0.19"
lambda_runtime = "0.6"
aws_lambda_events = "0.5"
futures = "0.3"
tokio = { version = "1", features=["full"] }
serde = { version = "1", features=["derive"] }
serde_json = "1"
```


A Rust Equivalent

```
[dependencies]
anyhow = "1"
async-tar = "0.4"
aws-config = "0.49"
aws-sdk-s3 = "0.19"
lambda_runtime = "0.6"
aws_lambda_events = "0.5"
futures = "0.3"
tokio = { version = "1", features=["full"] }
serde = { version = "1", features=["derive"] }
serde_json = "1"
```

← Need a different tar crate...

A Rust Equivalent

↓ ...because it's all going to be async

```
async fn index_tarball(  
    client: &s3::Client,  
    bucket: &str,  
    input_key: &str,  
    output_key: &str,  
) -> Result<()> {  
  
    // TODO: asyncify the previous code  
  
}
```

A Rust Equivalent

```
let tarball = async_tar::Archive::new(  
    client  
        .get_object()  
        .bucket(bucket)  
        .key(input_key)  
        .send()  
        .await?  
        .body  
        .map_err(|e| std::io::Error::new(std::io::ErrorKind::Other, e))  
        .into_async_read(),  
);
```

← make it impl `AsyncRead`

A Rust Equivalent

```
let mut output = Vec::new();  
let mut entries = tarball.entries()?;  
while let Some(entry) = entries.try_next().await? {  
    serde_json::to_writer(  
        &mut output,  
        &IndexEntry {  
            // unchanged from previous version  
        },  
    )?;  
    writeln!(output)?;  
}
```

← This is now a `Stream`,
not an `Iterator`

A Rust Equivalent

```
let output = tarball
    .entries()?
    .map_err(anyhow::Error::from)
    .try_fold(Vec::new(), |mut output, entry| async move {
        serde_json::to_writer(
            &mut output,
            &IndexEntry {
                // unchanged from previous version
            },
        )?;
        writeln!(output)?;
        Ok(output)
    })
    .await?;
```

← `TryStreamExt` has a lot of powerful helper methods

A Rust Equivalent

```
client
    .put_object()
    .bucket(bucket)
    .key(output_key)
    .body(output.into())
    .send()
    .await?;
```

A Rust Lambda

```
use cobalt_aws::lambda::{run_message_handler, Error};

#[tokio::main]
async fn main() -> Result<(), Error> {
    run_message_handler(message_handler).await
}

// Not shown: impls to populate this from env vars
struct Context {
    client: s3::Client,
    bucket: String,
}

async fn message_handler(input_key: String, context: Arc<Context>) -> Result<(),> {
    let output_key = ... // Not shown, for brevity
    index_tarball(&context.client, &context.bucket, input_key, &output_key)
        .await?;
}
```

Deploy via Docker Image

```
RUN mkdir /bin
```

```
RUN --mount=type=cache,target=/usr/local/cargo/registry \  
    --mount=type=cache,target=/build/target \  
    cargo build --profile release --target x86_64-unknown-linux-musl
```

```
RUN mv target/x86_64-unknown-linux-musl/index-tarballs-lambda /bin
```

```
ENTRYPOINT ["/bin/index-tarballs-lambda"]
```


Deploy via Docker Image

```
RUN mkdir /bin
```

```
RUN --mount=type=cache,target=/usr/local/cargo/registry \  
  --mount=type=cache,target=/build/target \  
  cargo build --profile release --target aarch64-unknown-linux-musl
```

Amazon suggest that this arch is usually cheaper



```
RUN mv target/x86_64-unknown-linux-musl/index-tarballs-lambda /bin
```

```
ENTRYPOINT ["/bin/index-tarballs-lambda"]
```

Rough Performance Numbers

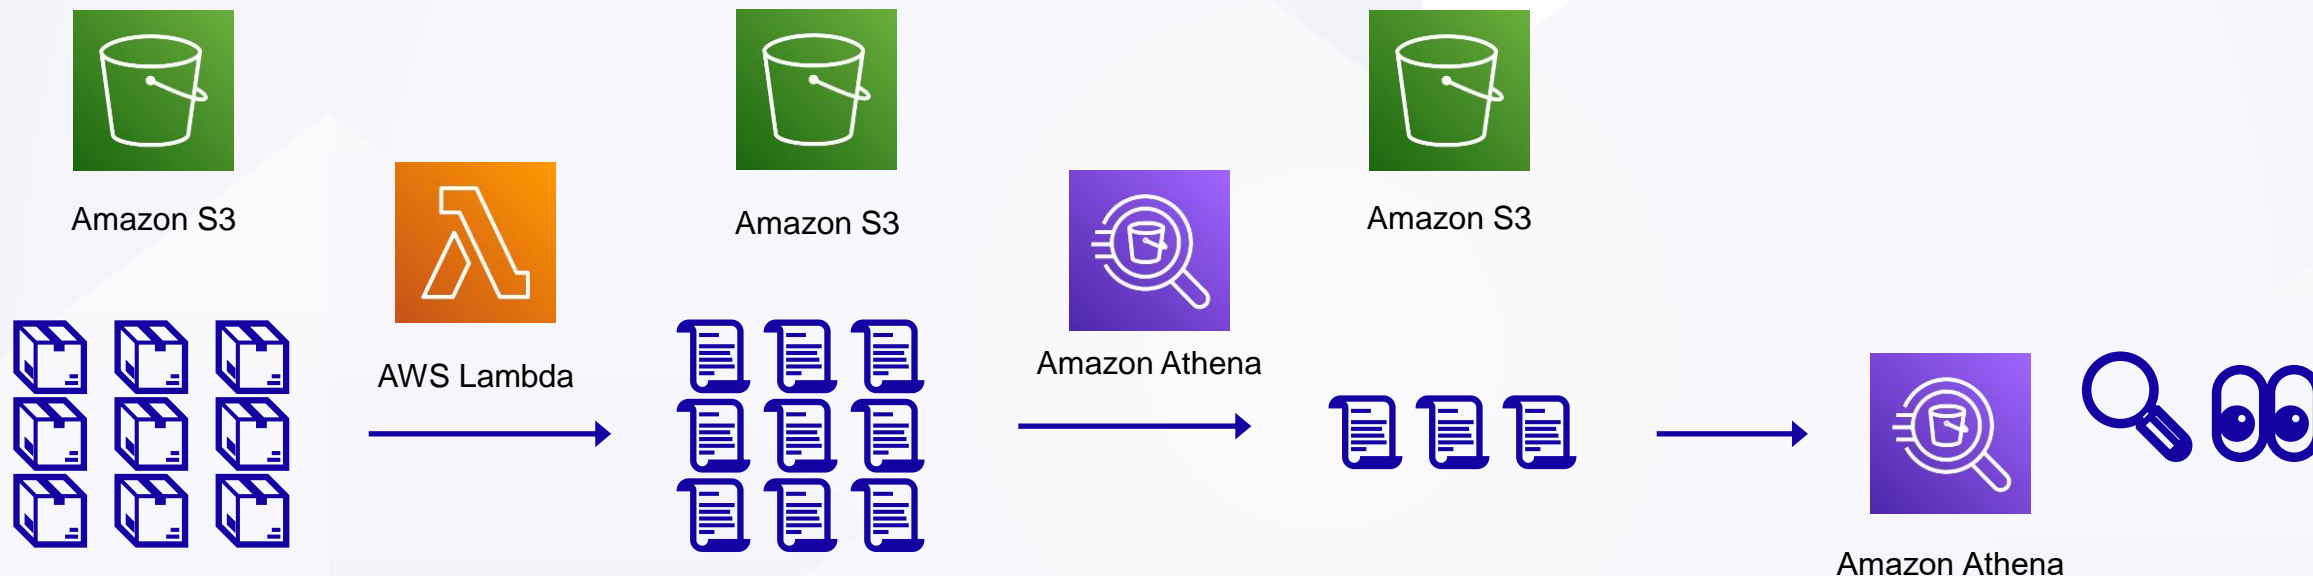
	architecture = x86_64	architecture = arm64
Python	14.5 seconds avg	14.3 seconds avg
Rust	7.3 seconds avg	7.8 seconds avg

Rough Performance Numbers

	architecture = x86_64	architecture = arm64
Python	\$241 / mil	\$190 / mil
Rust	\$121 / mil	\$104 / mil

The IRL Version...

- Generated 6 different listings per tarball, with rich metadata
- Partitioned output files by prefix, to help Athena
- Used transparent zstd compression on input and output files
 - (thanks, ``AsyncRead`/`AsyncWrite`` traits!)
- Cost $O(\$100)$ in lambda execution time



A few PB of de-identified
medical image data in
millions tar archives

.jsonl files listing
contents of each
archive

.parquet files
summarizing
available data



A few PB of de-identified
medical image data in
millions tar archives

.jsonl files listing
contents of each
archive

.parquet files
summarizing
available data



A few PB of de-identified medical image data in millions tar archives

.jsonl files listing contents of each archive

.parquet files summarizing available data

Reflections

Things we liked

- Runtime performance
- Low, stable memory usage
- Runtime robustness
 - “If it compiles, it works!”
- High-level abstractions
- Powerful async helpers
- Ease of build/packaging

Things that were challenging

- Async ecosystem fragmentation
- Testing/mockings
 - plug: LocalStack
- Runtime debugging context
 - plug: tracing
- Optimisation opportunities are an attractive nuisance

So, would we do it again?

We're making this a core competency

And trying to open-source where we can

- Higher-level abstractions for working with AWS
 - <https://github.com/harrison-ai/cobalt-aws/>
- Docker-based build tooling
 - <https://github.com/harrison-ai/dataeng-tooling-rust/>

thank you.

Ryan Kelly

<https://rfk.id.au/>



harrison.ai